

## Análisis y Reporte sobre el desempeño del modelo

1. Implementación elegida: uso de framework, sklearn (linear\_model)
  - a. Separación y evaluación del modelo con un conjunto de prueba y un conjunto de validación (Train/Test/Validation).

```
[17] # Score the model
reg.score(x_test, y_test)

0.8513618855525324

[18] reg.score(x_train, y_train)

0.8561208934839917

scores = cross_val_score(reg, x_train, y_train, cv=5, scoring='r2')
print("Mean score of %.2f with a standard deviation of %.2f" % (scores.mean(), scores.std()))

Mean score of 0.86 with a standard deviation of 0.01
```

- b. Diagnóstico y explicación el grado de bias y varianza: bajo medio alto

```
avg_expected_loss, avg_bias, avg_var = bias_variance_decomp(reg, x_train, y_train, x_test, y_test, loss='mse', num_rounds=50, random_seed=20)

# summary of the results
print('Average expected loss: %.3f' % avg_expected_loss)
print('Average bias: %.3f' % avg_bias)
print('Average variance: %.3f' % avg_var)
```

Al tratarse de una regresión lineal el bias es alto debido a que el modelo es poco flexible al ajuste de los datos, por otro lado, tiene una varianza baja (cambiar los datos de entrenamiento produce cambios pequeños en la estimación).

- c. Diagnóstico y explicación el nivel de ajuste del modelo: underfitt fitt overfitt

```
[76] # Model Information
print("Intercept:", reg.intercept_)
print("Coeficiente:", list(zip(x_test.columns, reg.coef_.flatten(), )))
print("Coeficiente de determinación R^2:", reg.score(x_test, y_test))

Intercept: [-351.02587529]
Coeficiente: [('height', 7.7225058871406205)]
Coeficiente de determinación R^2: 0.8513618855525324

# Model test error
predicciones = reg.predict(X = x_test)
rmse = mean_squared_error(y_true = y_test, y_pred = predicciones, squared = False)
print(f"El error (rmse) de test es: {rmse}")

El error (rmse) de test es: 12.372434423109533
```

2. Técnicas de regularización implementadas: Lasso y Ridge

a. Resultados:

```
[67] # Lasso model regularization
lasso_reg = Lasso(alpha=0.3)
lasso_reg.fit(x_train, y_train)

Lasso(alpha=0.3)

[68] lasso_reg.score(x_test, y_test)

0.8513712187225047

[69] lasso_reg.score(x_train, y_train)

0.8561149964826482

[70] # Ridge model regularization
ridge_reg = Ridge(alpha=0.3)
ridge_reg.fit(x_train, y_train)

Ridge(alpha=0.3)

[71] ridge_reg.score(x_test, y_test)

0.8513619002560128

[72] ridge_reg.score(x_train, y_train)

0.8561208934784967

[73] # Lasso and ridge predictions
prediction_lasso=lasso_reg.predict(x_test)
prediction_ridge=ridge_reg.predict(x_test)

[74] prediction_lasso=prediction_lasso.reshape(2000,1)
sns.distplot(y_test-prediction_lasso)

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:100: FutureWarning:
<matplotlib.axes._subplots.AxesSubplot at 0x7f15b98ab...>

Density
0.030
0.025
0.020
0.015
0.010
0.005
0.000
-40 -20 0 20 40
```

```
sns.distplot(y_test-prediction_ridge)

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:100: FutureWarning:
<matplotlib.axes._subplots.AxesSubplot at 0x7f15b9774550>

Density
0.030
0.025
0.020
0.015
0.010
0.005
0.000
-40 -20 0 20 40
```